
Triarch in Plain English

Bob Corbin — what I’m building, why it matters, and the part everybody seems to miss

The whole thing in one sentence:

I build the governance structure that lets regulated industries actually use AI without ending up in front of a regulator they can’t defend the decision to.

That’s the pitch. The rest of this is unpacking it.

What everybody gets wrong

Every conversation about AI right now is about how capable the models are getting. GPT-5 does this. Claude does that. Soon AI will do everything.

That’s true and it misses the point.

A more powerful engine doesn’t put more cars on public roads. What puts cars on public roads is *the rest of the car* — the brakes, the steering, the mirrors, the inspection sticker, the title, the insurance. None of which make the engine stronger. All of which are required to drive it legally.

Same with AI. The frontier AI companies are building amazing engines. But almost nobody is building the rest of the car for AI in regulated environments. Which is why hospitals, factories, banks, and utilities are stuck — they can see the engine, they want the engine, and they can’t legally use it the way it currently ships.

Why “regulated” is the whole game

Take a hospital trying to use AI for clinical decision support. AI suggests something. Patient outcome happens. Lawyer shows up. The lawyer asks: who decided this, on what basis, and where’s the audit trail?

If the answer is *“the AI suggested it, the doctor reviewed and accepted it, here are the signed records,”* that’s defensible. If the answer is *“the AI did it, somehow,”* that’s a settlement.

Take a Tier 1 auto supplier. AI helps draft a PPAP submission. OEM auditor asks: show me the change control, show me the engineer who signed off, show me the verification record.

If the AI output is wrapped in proper change control, signed by an accountable engineer, with full verification logs — defensible. If it’s a chatbot transcript — that’s a corrective action and possibly a loss of supplier status.

The pattern repeats across every regulated industry. Pharma. Aerospace. Energy. Finance. Banking. *The AI capability is not the bottleneck. The evidentiary structure around the AI is the bottleneck.* And nobody currently sells that structure to the people who need it.

That structure is what I build.

Why this isn’t already done

The frontier AI companies are mostly ML researchers. They optimize models, publish papers, ship developer APIs. They’re extraordinary at what they do.

What they are not is people who spent ten years on the floor of an automotive supplier learning what an IATF 16949 auditor actually checks. Or in a hospital learning what a Joint Commission survey actually requires. Or in a wastewater plant learning what a state DEP needs in your monthly discharge monitoring report.

That knowledge lives in industry. It lives in people who built compliance systems the hard way, got beat up by audits, and fixed the corrective actions afterwards. It doesn’t transfer to a research lab.

So you have AI companies racing to make their engines more powerful, regulated industries desperate to use those engines, and almost nobody in between making it possible.

That’s the gap. That’s where Triarch lives.

Why me

Ten years in manufacturing engineering. PPAP, IATF 16949, robot integration, PLC programming, quality documentation. I sat through the audits. I wrote the corrective actions. I learned what a defensible answer looks like because I had to produce one in front of auditors who’d seen every dodge.

When the modern AI wave arrived, I saw what was coming: people would try to deploy it in regulated environments and fail, because nobody was building the governance layer. So I started building it.

Triarch is the company doing that build. The principles are imported from manufacturing — every claim auditable, every decision reversible, every system designed to survive its own failure modes rather than hope they don’t happen. The output is software and structured process, not hardware. The customers are organizations that need AI capability with regulator-defensible discipline.

The bet

As AI gets more capable, demand for *governance around AI* grows faster than demand for AI itself. Because the consequences of ungoverned AI in a regulated environment scale with the AI's capability. A weak AI that hallucinates costs you an embarrassing email. A powerful AI that hallucinates costs you a federal investigation.

The frontier labs win the substrate. Vertical specialists with real domain depth win the governance layer. The market for both grows. They're complements, not competitors.

I'm betting on the governance layer, in regulated manufacturing first, then adjacent domains. That's the work.

What I want from this conversation

If you know someone in manufacturing, healthcare, finance, or any heavily-regulated industry who's been told "*use AI, but make it auditable*" and is currently failing at that — that's my customer. Introduce me.

If you don't, but the idea makes sense to you and you want to argue with it — also great. The framework gets sharper every time someone pushes on it.

What I'm not looking for: another conversation about how powerful the models are getting. I know. The capability isn't the question. The governance is.